

LINFO1104 – LSINC1104
Concepts, paradigms, and semantics
of programming languages

Lectures 8 & 9
Exception handling
and concurrent programming

Peter Van Roy

ICTEAM Institute
Université catholique de Louvain

peter.vanroy@uclouvain.be



1

Overview of lectures 8-9

- Exception handling
 - How to handle exceptional situations in a program without making the program more complicated
- Concurrent programming
 - Deterministic dataflow (a.k.a. functional dataflow)
 - Semantics of concurrent programming
 - “Concurrency for Dummies”
 - Programming techniques for deterministic dataflow



2

Exception handling



3

How to handle exceptional situations



- How can we handle exceptional situations in a program?
 - Such as: division by 0, opening a nonexistent file, and so forth
 - Program errors but also errors from outside the program
 - Things that happen rarely but that must be taken care of
- We add a new programming concept called **exceptions**
 - We define exceptions and show how they are used
 - We give the semantics of exceptions in the abstract machine
- With exceptions, we can handle exceptional situations without cluttering up the program with rarely used error checking code

4


The containment principle



- When an error occurs in a program, we would like to be able to recover from the error
 - The program should not crash
- Furthermore, we would like the error to affect as little as possible of the program
- We propose **the containment principle**:
 - A program is a set of **nested execution contexts**
 - An error will occur **inside** an execution context
 - A recovery routine (exception handler) exists at the boundary of an execution context, to make sure the error **does not propagate** to higher execution contexts

5

Handling an exception



- ✗ An error that raises an exception
- An execution context
- The execution context that catches the exception

- An executing program that encounters an error will jump to another part (the exception handler) with an argument (the exception) that describes the error
 - The exception handler fixes incorrect data structures so the program can continue

6

The try and raise instructions

- We introduce two new instructions for handling exceptions:

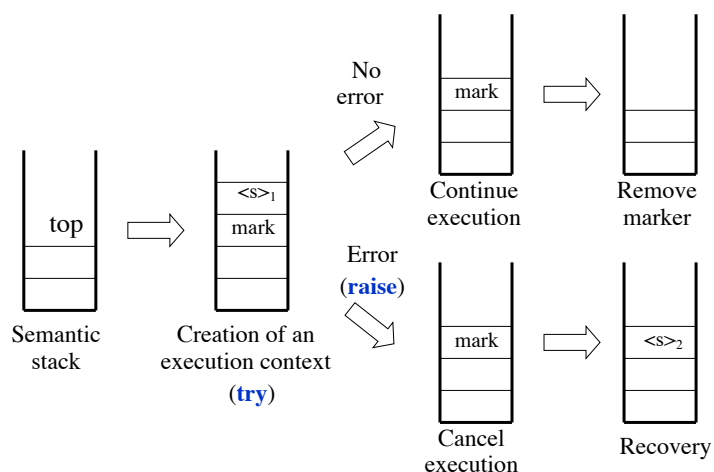
```
try <s>1 catch <y> then <s>2 end % Create an execution context
raise <x> end % Raise an exception
```

- With the following behavior:

- try** puts a “marker” on the stack and starts executing <s>₁
- If there is no error, <s>₁ executes normally and removes the marker when it terminates
- raise** is executed when there is an error, which empties the stack up to the marker (the rest of <s>₁ is therefore canceled)
 - Then <s>₂ is executed
 - <y> refers to the same variable as <x>
 - The scope of <y> exactly covers <s>₂

7



Semantics of exceptions



8

An execution context

- An **execution context** is the part of the semantic stack that starts with a marker and continues to the stack top:



9

Example using exceptions

```

fun {Eval E}
  if {IsNumber E} then E
  else
    case E
    of plus(X Y) then {Eval X}+{Eval Y}
    [] times(X Y) then {Eval X}*{Eval Y}
    else raise badExpression(E) end
    end
  end
end

try
  {Browse {Eval plus(23 times(5 5))}}
  {Browse {Eval plus(23 minus(4 3))}}
catch X then {Browse X} end
  
```

- Using exceptions, the error handling code does not clutter up the program

10

If we did not have exceptions...



```
fun {Eval E}
  if {IsNumber E} then E
  else
    case E
    of plus(X Y) then R={Eval X} in
      case R of badExpression(RE) then badExpression(RE)
      else R2={Eval Y} in
        case R2 of badExpression(RE) then badExpression(RE)
        else R+R2
        end
      end
    [] times(X Y) then
      % ... Same code as plus
    else badExpression(E)
    end
  end
end
```

- Much more code!
 - In this example, 22 lines instead of 10 (more than double)
- The code is much more complicated because of all the **case** statements handling badExpression

11

The “finally” clause



- The **try** has an additional **finally** clause, for an operation that must always be executed (in both the correct and error cases):

```
FH={OpenFile “foobar”}
try
  {ProcessFile FH}
catch X then
  {Show “*** Exception during execution ***”}
finally {CloseFile FH} end % Always close the file
```

12

Exceptions in Java



- An exception is an object that inherits from the class Exception (which is a subclass of Throwable)
- There are two kinds of exceptions
 - **Checked exceptions:** The compiler verifies that all methods only throw the exceptions declared for the class
 - **Unchecked exceptions:** Some exceptions can arrive without the compiler being able to verify them. They inherit from RuntimeException and Error.
- For exceptions that the program itself defines, you should **always use checked exceptions**, since they are declared and therefore part of the program's interface

13

Java exception syntax



```
throw new NoSuchElementException(name) ;

try {
    <stmt>
} catch (exctype1 id1) {
    <stmt>
} catch (exctype2 id2) {
    ...
} finally {
    <stmt>
}
```

14



Good style

- We read a file and perform an action for each item in the file:

```
try
    while (!stream.eof())
        process(stream.nextTokn());
finally
    stream.close();
```

15



Bad style

- We can use the exception handler to change the execution order **during normal execution**:

```
try {
    for (;;)
        process (stream.next());
} catch (StreamEndException e) {
    stream.close();
}
```

- Reaching the end of a stream is completely **normal**, it is not an error. What happens if a **real error** happens and is mixed in with the normal operation? You don't want to mix things. Normal operation should be kept separate from errors!

16

Conclusions on data abstraction



17

Java, Scala, and language design



- Java was designed to support data abstraction (1990s)
 - True data abstraction (encapsulation, GC)
 - All entities are objects or ADTs
 - Support for object-oriented design principles
- Scala has added two principles to this (2000s)
 - Separation between mutable/immutable (functional programming)
 - Everything is an object (including functions)
- These principles considerably increase Scala's expressive power compared to Java
 - We consider that Scala is an important successor to Java
 - Although some people consider it is a Swiss Army knife!

18

Final remarks



- This completes the part of the course on data abstraction
 - Explicit state (cells), objects, and ADTs
 - Exceptions
 - We did not go in-depth into object-oriented programming techniques because they are covered in other courses
- So far in the course we have covered **three important themes**
 - **Functional programming** (including recursion, invariant programming, and higher-order programming)
 - **Language semantics** (including a complete operational semantics and an introduction to lambda calculus)
 - **Data abstraction** (including explicit state, objects, and ADTs)
- The next theme is **concurrent programming**

19

Concurrent programming



20

The world is concurrent



- The real world is **concurrent**
 - It is made of **activities that progress independently**
- The computing world is concurrent too:
 - **Distributed system**: computers linked by a network
 - A concurrent activity is called a **computing node (computer)**
 - Each computing node has its own resources (memory, CPU)
 - **Operating system**: management of a single computer
 - A concurrent activity is called a **process**
 - Processes share the same computer resources and have independent memory spaces
 - **Process**: execution of a single program
 - A concurrent activity is called a **thread**
 - Threads share the same memory space

21

Concurrent programming



- Concurrency is natural
 - Many activities are naturally independent
 - Activities that are **independent** are ipso facto **concurrent**
 - So how can we write a program with many independent activities?
 - Concurrency must be supported by the language!
- A concurrent program
 - Multiple progressing activities that exist at the same time
 - Activities that can communicate and synchronize
 - **Synchronize**: an activity waits for an action of another activity
 - **Communicate**: information passes from one activity to another

22

Concurrency can be (very) hard



- It introduces many difficulties such as nondeterminism, race conditions, reentrancy, deadlocks, livelocks, fairness, consistency of shared data
 - Java's [synchronized objects \(monitors\)](#) are tough to program with
 - Erlang's and Scala's [actors](#) are better, but they still have race conditions
 - [Libraries](#) can hide some of these problems, but they always peek through
- Adding distribution (networked systems) makes it **even harder**
- Adding partial failure makes it **even much harder than that**
- The Holy Grail: can we make concurrent programming as easy as sequential programming?
 - Yes, it can be done, if the paradigm is chosen wisely
 - In this course we will see **deterministic dataflow**, which is a concurrent paradigm that is a form of functional programming

← INFO1131

INFO2345

23

Deterministic dataflow



24

Concurrency paradigms



- There are **three main paradigms** of concurrent programming
- **Deterministic dataflow** (the simplest and best)
 - This paradigm is also called functional dataflow
 - It supports all the techniques of functional programming
 - That is what we will see today
- What are the two other paradigms?
 - **Message-passing concurrency** (e.g., Erlang and Scala actors)
 - Activities send messages to each other (like sending letters)
 - This works well and is not too hard
 - **Shared-state concurrency** (e.g., Java monitors)
 - Activities share the same data and they try to work together without getting in each other's way
 - Much more complicated than the two previous paradigms
 - Unfortunately, many current languages still use this paradigm

We will see Erlang in the course

LINFO1131

25

An unbound variable



- Let us explain dataflow by starting with an unbound variable
- An unbound variable is created in memory but not bound to a value
- What happens when you invoke an operation with an unbound variable?
local X Y **in**
 Y=X+1
 {Browse Y}
end
- What happens?

26

What to do with an uninitialized variable?



- Different languages do different things
 - In **C**, the addition continues and X has a “garbage value” (= content of X's memory at that moment)
 - In **Java**, the addition continues and X's value is 0 (if X is an object attribute with type integer)
 - In **Prolog**, execution stops with an error
 - In **Java**, the compiler detects an error (if X is a local variable)
 - In **Oz**, execution waits just before the addition and continues when X is bound (dataflow execution)
 - In **constraint programming**, the equation “ $Y=X+1$ ” is added to the set of constraints and execution continues. An amazing way to compute!

LINFO2365
Programmation
par contraintes

27

Continuing the execution



- The waiting instruction:
declare X
local Y **in**
 $Y=X+1$
 {Browse Y}
end
- If someone would bind X, then execution could continue
- But who can do it?

28


Continuing the execution



- The waiting instruction:
`declare X`
`local Y in`
 $Y = X + 1$
 {Browse Y}
`end`
- If someone would bind X, then execution could continue
- But who can do it?
- Answer: another concurrent activity!
- If another activity does:
 $X = 20$
- Then the addition will continue and display 21!
- This is called **dataflow execution**

29

Dataflow execution



- Activity A waits patiently at point (1) just before the addition
- When activity B binds $X=20$ at point (2), then activity A can continue
- If activity B binds $X=20$ **before** activity A reaches point (1), then activity A **does not have to wait**

30

Threads



31

Threads



- We add a language concept to support concurrent activities
 - In a program, an activity is a **sequence of executing instructions**
 - We add this concept to the language and call it a **thread**
- Each thread is **sequential**
- Each thread runs **independently** of the others
 - There is no order defined between different threads
 - The system executes all threads using **interleaving semantics**: it is as if only one thread executes at a time, with execution switching rapidly from one thread to another
 - The system guarantees that each thread receives a fair share of the computational capacity of the processor
- Two threads can communicate if they share a variable
 - For example, the variable corresponding to identifier X in the example we just saw

32



Thread creation

- Creating a thread in Oz is simple
- Any instruction can be executed in a new thread:
`thread <s> end`
- For example:
`declare X`
`thread {Browse X+1} end`
`thread X=1 end`
- What does this small program do?
 - Several executions are possible, but they all eventually arrive at the same result: 2 is displayed!

33



A small program (1)

- A small program with several threads:
`declare X0 X1 X2 X3 in`
`thread X1=1+X0 end`
`thread X3=X1+X2 end`
`{Browse [X0 X1 X2 X3]}`
- The Browser displays [X0 X1 X2 X3]
 - The variables are all unbound
 - The Browser also uses dataflow:
when a variable is bound, the display is updated

34



A small program (2)

- A small program with several threads:
declare X0 X1 X2 X3 **in**
thread X1=1+X0 **end**
thread X3=X1+X2 **end**
{Browse [X0 X1 X2 X3]}
- Two threads will wait:
 - X1=1+X0 waits (since X0 is unbound)
 - X3=X1+X2 waits (since X1 and X2 are unbound)

35



A small program (3)

- A small program with several threads:
declare X0 X1 X2 X3 **in**
thread X1=1+X0 **end**
thread X3=X1+X2 **end**
{Browse [X0 X1 X2 X3]}
- Let's bind one variable
 - Bind X0=4

36



A small program (4)

- A small program with several threads:
declare X0 X1 X2 X3 **in**
thread X1=1+X0 **end**
thread X3=X1+X2 **end**
{Browse [X0 X1 X2 X3]}
- Let's bind one variable
 - Bind X0=4
 - The first thread executes and binds X1=5
 - The Browser displays [4 5 _ _]

37



A small program (5)

- A small program with several threads:
declare X0 X1 X2 X3 **in**
thread X1=1+X0 **end** % terminated
thread X3=X1+X2 **end**
{Browse [X0 X1 X2 X3]}
- The second thread is still waiting
 - Because X2 is still unbound

38

A small program (6)



- A small program with several threads:
`declare X0 X1 X2 X3 in`
`thread X1=1+X0 end % terminated`
`thread X3=X1+X2 end`
`{Browse [X0 X1 X2 X3]}`
- Let's do another binding
 - Bind $X2=7$
 - The second thread executes and binds $X3=12$
 - The Browser displays [4 5 7 12]

39

The Browser is a dataflow program



- The Browser executes with its own threads
- For each unbound variable that is displayed, there is a thread in the Browser that waits until the variable is bound
 - When the variable is bound, the display is updated
- This does not work with cells
 - The Browser uses the functional dataflow paradigm
 - The Browser does not look at the content of cells

40

Streams and agents



41

Streams



- A **stream** is defined as a list that ends in an unbound variable
 - $S = a|b|c|d|S2$
 - A stream can be extended with new elements as long as necessary
 - The stream can be closed by binding the end to nil
- A stream can be used as a **communication channel** between two threads
 - The first thread adds elements to the stream
 - The second thread reads the stream

42

Programming with streams



- This program displays the elements of a stream as they appear:

```
proc {Disp S}
  case S of X|S2 then {Browse X} {Disp S2} end
end
declare S
thread {Disp S} end
```

- We can add elements gradually:
 declare S2 **in** S=a|b|c|S2
 declare S3 **in** S2=d|e|f|S3
- Try it yourself!

43


Producer/consumer (1)



- A **producer** generates a stream of data
 fun {Prod N} {Delay 1000} N|{Prod N+1} **end**
 - The {Delay 1000} slows down execution enough to observe it
- A **consumer** reads the stream and performs some action (like the Disp procedure)
- A producer/consumer program:
 declare S
 thread S={Prod 1} **end**
 thread {Disp S} **end**

44

Producer/ consumer (2)



- Each circle is a concurrent activity that reads and writes streams
 - We call this an agent
- Agents P and C communicate through stream S
 - The first thread creates the stream, the second reads it

45

Pipeline (1)




- We can add more agents between P and C
- Here is a transformer that modifies the stream:

```
fun {Trans S}
  case S of X|S2 then X*X|{Trans S2} end
end
```
- This program has three agents:

```
declare S1 S2
thread S1={Prod 1} end
thread S2={Trans S1} end
thread {Disp S2} end
```

46

Pipeline (2)



- We now have three agents
 - The producer (agent P) creates stream S1
 - The transformer (agent T) reads S1 and creates S2
 - The consumer (agent C) reads S2
- The pipeline is a very useful technique!
 - For example, it is **omnipresent in operating systems since Unix**

47

Agents

- An agent is a concurrent activity that reads and writes streams
 - The simplest agent is **a list function executing in one thread**
 - Since list functions are tail-recursive, the agent can execute with a fixed memory size
 - This is **the deep reason why single assignment is important**: it allows tail-recursive list functions, which makes deterministic dataflow into a practical paradigm
- All list functions can be used as agents
 - All functional programming techniques can be used in deterministic dataflow
 - Including higher-order programming! Later on we will see more examples of the power of the model.

48

Thread semantics



49

Thread semantics (1)




- We extend the abstract machine with threads
- Each thread has one semantic stack
 - The instruction **thread** <s> **end** creates a new stack
 - All stacks share the same memory
- There is **one sequence of execution states**, and threads take turns executing instructions
 - $(MST_1, \sigma_1) \rightarrow (MST_2, \sigma_2) \rightarrow (MST_3, \sigma_3) \rightarrow \dots$
 - MST is a multiset of semantic stacks
 - Each step " \rightarrow " executes one step in one thread
 - The choice of which thread to execute is made by the **scheduler**
 - This is called **interleaving semantics**

50

Thread semantics (2)



A semantic stack that is about to create a thread




51

Thread semantics (3)



We now have two stacks!



52

Why interleaving semantics?



- **Interleaving semantics** is much easier to reason about than true concurrency semantics
 - **True concurrency semantics** = more than one thread can execute in one execution step
- Imagine that all threads execute in parallel, each with its own processor but all sharing the same memory
 - What happens when two threads write simultaneously at the same memory word?
 - With interleaving semantics, one thread will always write before the other, which makes reasoning simple
 - True concurrency semantics also models where threads "step on each others' toes", but usually this is not needed, since the hardware is designed so that this does not happen
 - For example, in a multicore processor the **cache coherence protocol** avoids simultaneous operations on one memory word

53

Concurrent program execution




54

Total order of a sequential program



- A sequential program is a program with **one thread**
- In a sequential program, execution states are in a **total order**
 - **Total order** = there is a defined order between all pairs of states




55

Partial order of a concurrent program

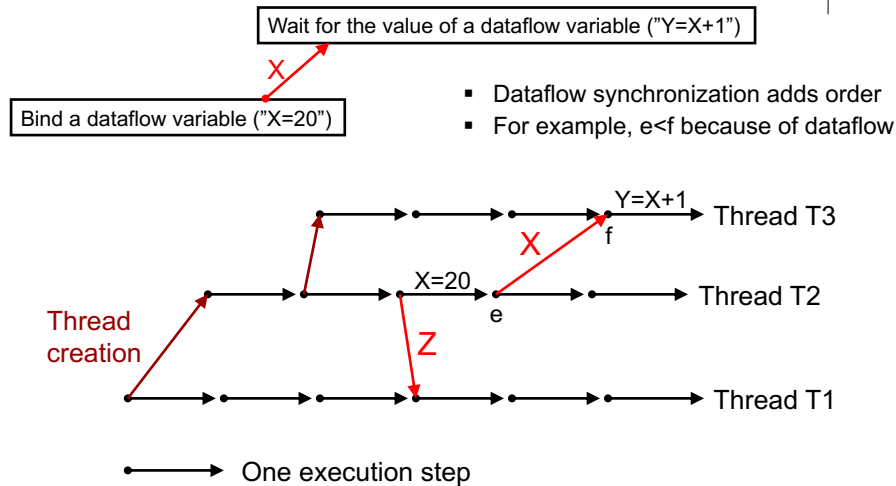


- A concurrent program is a program with **more than one thread**
- In a concurrent program, execution states are in a **partial order**
 - **Partial order** = not all pairs of execution states have a defined order
 - For example, $c < d$ (c before d) but b and c have no order



56

Partial order of a concurrent program




57

The actual execution order



- The processor execution is always one sequence of execution states
- The scheduler chooses the order of these states, which is always compatible with the partial order (choice = **nondeterminism**)



58

Nondeterminism



59

Nondeterminism and the scheduler



- **Nondeterminism** is the ability of a system to make decisions (choices) independently of the system's developer
- The **scheduler** is the part of the system that decides at each moment which thread to execute
 - This decision is an example of **nondeterminism**
 - Scheduler decisions often vary from one execution to the next; they depend on external conditions such as processor load, memory behavior (caching), network behavior, and timing of external events
- Nondeterminism exists in all concurrent systems
 - It must be so, since the concurrent activities are **independent**
 - All concurrent programs must manage their nondeterminism!

60

Example of nondeterminism (1)



- What does the following program do?

```
declare X
thread X=1 end
thread X=2 end
```
- The execution order of the two threads is not fixed
 - X will be bound to 1 or 2, we don't know which
 - The other thread **will have an error (raise an exception)**
 - A variable cannot be assigned to two values
- This is an example of **nondeterminism**
 - **A choice made by the system during execution**
 - The system is free to choose one or the other

61

Example of nondeterminism (2)



- What does the following program do?

```
declare X={NewCell 0}
thread X:=1 end
thread X:=2 end
```
- The execution order of the two threads is not fixed
 - Cell X will first be bound to one value, then to the other
 - When both threads terminate, X will contain 1 or 2, we don't know which
 - This time there is no error
- This is also an example of **nondeterminism**
 - **A choice made by the system during execution**

62

Example of nondeterminism (3)



- What does the following program do?

```
declare X={NewCell 0}  
thread X:=1 end  
thread X:=1 end
```
- It makes a choice, just like the previous program
 - But in this case, the final results are the same (by accident)
- **This is still nondeterminism!**
 - The important point is the **choice**: the running program still sees a difference in the threads' execution order
 - The results may be the same by accident (depending on the computations done), but the choice remains

63

Managing nondeterminism



- Nondeterminism *must always be managed*
 - It should not affect program correctness (this can be very tricky!)
 - The most complicated case is **threads and cells used together** (see previous example)
 - Unfortunately, this is exactly how many languages handle concurrency (Java, C++, C#, etc.) → see course LINFO1131
- Deterministic dataflow has a major advantage
 - **The result of a program is always the same** (except if there is a program error, that is, if a thread raises an exception)
 - The nondeterminism of the scheduler **does not affect the result**
 - There is no **observable** nondeterminism
 - We call this « Concurrency for Dummies »
 - It is a consequence of Church-Rosser (functional programming)

64



How the scheduler works (1)

- The choice of which thread to execute and for how long is made by the scheduler
- Time slices (on modern systems, often 10ms)
 - Each thread executes during a short time period called a **time slice**
 - On multicore processors, some operating systems can allow time slices on different cores, but there is still interleaving semantics
- Thread states (runnable and suspended)
 - A thread is **runnable** if the instruction on the top of its stack is not waiting on a dataflow variable. Otherwise, the thread is **suspended**, in other words **blocked on a variable**.

65



How the scheduler works (2)

- Fairness
 - A scheduler is **fair** if every runnable thread will eventually (eventually = in finite time) be executed
 - Usually, threads are classified according to their **priority**, and some **additional guarantees** are given on the percentage of the processor time that is given to the threads of the same priority
- Reasoning about programs
 - If the scheduler is fair, then it is possible to reason about program execution (all threads will run to completion)
 - If the scheduler is not fair, then a perfectly correct program may not run correctly
 - Certain threads may **starve**, i.e., receive 0% of the processor time, so they never execute, and the program just stops

66

“Concurrency for dummies”



67

“Concurrency for dummies”



- The multi-agent programs we saw so far are all **deterministic**
 - Their nondeterminism is not observable (results always the same)
 - The agent Trans with input 1|2|3|_ always outputs 1|4|9|_
- In these programs, concurrency does not change the result but only **the order in which computations are done** (that is, **when** the result is calculated may be different)
 - This is **always true** in deterministic dataflow
 - It is possible to add threads at will to a program without changing the result (we call this **Concurrency for Dummies**)
 - The only effect of added threads is to make the program more incremental (by interleaving execution and removing deadlocks)
- This only works in **functional programming** (deterministic dataflow)!
 - It is a consequence of the Church-Rosser theorem
 - It is not true when using cells and threads together (Java!)

68



Example (1)

```
fun {Map Xs F}
  case Xs
  of nil then nil
  [] X|Xr then
    {F X} | {Map Xr F}
  end
end
```

69



Example (2)

```
fun {CMap Xs F}
  case Xs
  of nil then nil
  [] X|Xr then
    thread {F X} end | {CMap Xr F}
  end
end
```

70



Example (3)

```
fun {CMap Xs F}
  case Xs
  of nil then nil
  [] X|Xr then
    thread {F X} end | {CMap Xr F}
  end
end
```

thread ... end
can be used as
an expression

71



Example (4)

```
fun {CMap Xs F}
  case Xs
  of nil then nil
  [] X|Xr then
    thread {F X} end | {CMap Xr F}
  end
end
```

- What happens when we execute:
declare F
{Browse {CMap [1 2 3 4] F}}

72



Example (5)

```
fun {CMap Xs F}
  case Xs
  of nil then nil
  [] X|Xr then
    thread {F X} end | {CMap Xr F}
  end
end
```

```
declare F
{Browse {CMap [1 2 3 4] F}}
```

- The Browser displays [_ _ _ _]
 - CMap calculates a list with unbound variables
 - The new threads wait until F is bound
- What would happen if {F X} was not in its own thread?
 - Nothing would be displayed! The CMap call would block.

73



Example (6)

```
fun {CMap Xs F}
  case Xs
  of nil then nil
  [] X|Xr then
    thread {F X} end | {CMap Xr F}
  end
end
```

- What happens when we bind F:
F = fun {\$ X} X+1 end

74



Example (7)

```
fun {CMap Xs F}
  case Xs
  of nil then nil
  [] X|Xr then
    thread {F X} end | {CMap Xr F}
  end
end
```

- The Browser displays [2 3 4 5]
- With or without the thread creation, the final result is always [2 3 4 5]

75



Concurrency for dummies!

- Threads can be added at will to a functional program **without changing the result**
- Therefore it is very easy to take a functional program and make it concurrent
- It suffices to insert **thread ... end** in those places that need concurrency
- **Warning:** concurrency for dummies does not work in a program with cells (= mutable variables)!
 - For example, it does not work in Java
 - In Java, concurrency is handled with the concept of a **monitor** (= **synchronized object**), which coordinates how multiple threads access an object. This is *much more complicated* than deterministic dataflow.

76

Why does it work? (1)



```
fun {Fib X}
  if X==0 then 0
  elseif X==1 then 1
  else
    thread {Fib X-1} end + {Fib X-2}
  end
end
```

77

Why does it work? (2)




```
fun {Fib X}
  if X==0 then 0 elseif X==1 then 1
  else F1 F2 in
    F1 = thread {Fib X-1} end
    F2 = {Fib X-2}
    F1 + F2
  end
end
```

Dataflow dependency

It works because variables can only be bound to one value (single assignment)

78

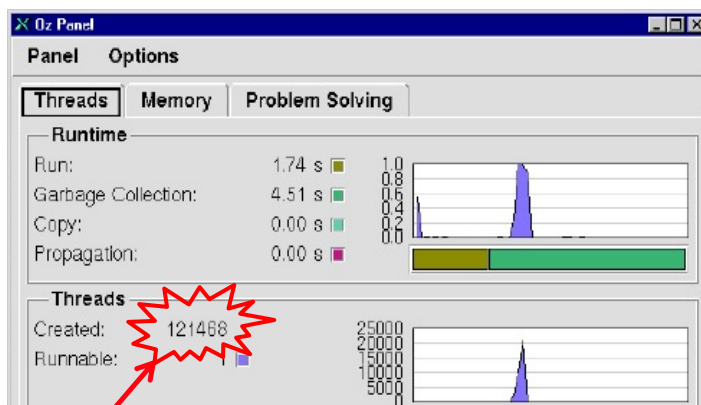
Execution of {Fib 6}



79

Observing the execution of Fib

Only in Mozart 1



Total number of threads created since system startup

Oz Compiler Panel (in Oz menu)

80

Counting threads

C={NewCell 0}

proc {Inc C}

{Exchange C X Y} Y=X+1

end

fun {Fib X}

if X==0 **then** 0

elseif X==1 **then** 1

else

thread {Inc C} {Fib X-1} **end** + {Fib X-2}

end

end

This works also in Mozart 2

C:= @C + 1 is not correct!
It is because the scheduler can put
the thread to sleep after the X=@C+1
and before the C:=X.

81

Multi-agent programming

82


Multi-agent programming



- Earlier in the course we saw some simple examples of multi-agent programs
 - Producer/consumer
 - Producer/transformer/consumer (pipeline)
- Let's see two more sophisticated examples
 - **Sieve of Eratosthenes**: dynamically building a pipeline during its execution
 - **Digital logic simulation**: using higher-order programming together with concurrency

83

The Sieve of Eratosthenes



- The Sieve of Eratosthenes is an algorithm for calculating a sequence of prime numbers
- Each agent in the pipeline removes multiples of an integer
- Starting with a sequence containing all integers, we end up with a sequence of primes

84



A filter agent

- A list function that removes multiples of K:

```
fun {Filter Xs K}
  case Xs of X|Xr then
    if X mod K \= 0 then X|{Filter Xr K}
    else {Filter Xr K} end
  else nil
  end
end
```

- We make an agent by putting it in a thread:

```
thread Ys={Filter Xs K} end
```

85



The Sieve program

- Sieve builds the pipeline during execution:

```
fun {Sieve Xs}
  case Xs
  of nil then nil
  [] X|Xr then X|{Sieve thread {Filter Xr X} end}
  end
end

declare Xs Ys in
  thread Xs={Prod 2} end
  thread Ys={Sieve Xs} end
{Browse Ys}
```

Concurrent deployment

Building the infrastructure of a concurrent program during its execution (execution will just wait if a part that it needs is not built yet)

86

An optimization



- Otherwise too many do-nothing agents are created!

```
fun {Sieve2 Xs M}  
  case Xs  
  of nil then nil  
  [] X|Xr then  
    if X=<M then  
      X|{Sieve2 thread {Filter Xr X} end M}  
    else Xs end  
  end  
end
```

- We call {Sieve2 Xs 316} to generate a list of primes up to 100000 (why?)

87

Digital logic simulation



88

Digital logic simulation



- The deterministic dataflow paradigm makes it easy to model digital logic circuits
- We show how to model combinational logic circuits (no memory) and sequential logic circuits (with memory)
- Signals in time are represented as streams; logic gates are represented as agents

89

Modeling digital circuits



- Real digital circuits consist of active circuit elements called gates which are interconnected using wires that carry digital signals
- A **digital signal** is a voltage in function of time
 - Digital signals are meant to carry two possible values, called 0 and 1, but they may have noise, glitches, ringing, and other undesirable effects
- A **digital gate** has input and output signals
 - The output signal is slightly delayed with respect to the input
- We will model **gates as agents** and **signals as streams**
 - This assumes perfectly clean signals and zero gate delay
 - We will later add a delay gate in order to model gate delay

90

Digital signals as streams




- A signal is modeled by a stream that contains elements with values 0 or 1

$$S = a_0 | a_1 | a_2 | \dots | a_i | \dots$$

- Time instants are numbered from when the circuit starts running
 - This models a clocked circuit
- At instant i , the signal's value $a_i \in \{0, 1\}$

91

Digital logic gates



x	y	z			
		Not	And	Or	Xor
0	0	1	0	0	0
0	1	1	0	1	1
1	0	0	0	1	1
1	1	0	1	1	0

- Some typical logic gates with their standard pictorial symbols and the boolean functions that define them
- But gates are not just boolean functions!

92

Digital gates as agents




- A gate is much more than a boolean function; it is an **active entity** that takes input streams and calculates an output stream

```
fun {And A B} if A==1 andthen B==1 then 1 else 0 end end
fun {Loop S1 S2}
  case S1#S2 of (A|T1)#(B|T2) then {And A B}|{Loop T1 T2} end
end
thread Sc={Loop Sa Sb} end
```

- Example execution:

```
Sx=0|1|0|Tx % input signal x
Sy=1|1|0|Ty % input signal y
Sz=0|1|0|Tz % output signal z
```



93

Creating many gates



- Let us define an **abstraction** for building all the different kinds of logic gates we need
 - We define the function GateMaker that takes a two-argument boolean function Fun, where {GateMaker Fun} returns a function FunG that creates gates
 - Each call to FunG creates a running gate based on Fun
- This gives **three levels of abstraction** that we can compare with object-oriented programming:
 - GateMaker is analogous to a **generic class** or **metaclass**
 - FunG is analogous to a **class**
 - A running gate is analogous to an **object**

94

GateMaker implementation



- Calling {GateMaker F} creates a gate maker:

```
fun {GateMaker F}
  fun {$ Xs Ys}
    fun {GateLoop Xs Ys}
      case Xs#Ys of (X|Xr)#(Y|Yr) then
        {F X Y}{{GateLoop Xr Yr}}
      end
    end
  in
    thread {GateLoop Xs Ys} end
  end
end
```

95

Making gates



- Each of these functions can make gates:

```
AndG={GateMaker fun {$ X Y} X*Y end}
OrG={GateMaker fun {$ X Y} X+Y-X*Y end}
NandG={GateMaker fun {$ X Y} 1-X*Y end}
NorG={GateMaker fun {$ X Y} 1-X-Y+X*Y end}
XorG={GateMaker fun {$ X Y} X+Y-2*X*Y end}
```

96

Combinational logic



97

Combinational logic




- Combinational logic has no memory: all calculation is done at the same time instant
- A gate is a simple combinational function:



- Therefore, any number of interconnected gates also defines a combinational function
- We define a useful circuit called a **full adder**

98

Full adder specification



x	y	z	c	s
0	0	0	0	0
0	0	1	0	1
0	1	0	0	1
0	1	1	1	0
1	0	0	0	1
1	0	1	1	0
1	1	0	1	0
1	1	1	1	1

- A full adder adds three 1-bit binary numbers x , y , and z giving a sum bit s and carry bit c
- An n -bit adder can be built by connecting n full adders

99

Full adder implementation



- Full adder creation as five-argument component:

```

proc {FullAdder X Y Z C S}
  A B D E F
in
  A={AndG X Y}
  B={AndG Y Z}
  D={AndG X Z}
  F={OrG B D}
  C={OrG A F}
  E={XorG X Y}
  S={XorG Z E}
end
  
```

100

Sequential logic



101

Sequential logic




- Sequential logic has memory: past values of a signal influence the present values
- We add a way for the past to influence the present: a Delay gate

$S = a_0|a_1|a_2|\dots|a_i|\dots$

$T = b_0|b_1|b_2|\dots|b_i|\dots$


$b_i = a_{i-1} \Rightarrow T = 0|S$
(if $i > 0$)



fun {DelayG S} 0|S **end**

102

Latch specification



- A latch is a simple circuit with memory; it has two stable states and can memorize its input
- Output d_o follows input d_i and freezes when c is 1

103

Latch implementation

- Latch creation as a three-argument component:

```

proc {Latch C Di Do}
  A B E F
in
  F={DelayG Do}
  A={AndG C F}
  E={NotG C}
  B={AndG E Di}
  Do={OrG A B}
end
  
```

104

Summary and history



105

Deterministic dataflow summary



- We have introduced a simple and expressive paradigm for concurrent programming
 - We can build multi-agent programs using **streams** (list with unbound tail) and **agents** (list function running in a thread)
- It is based on two simple ideas
 - **Single-assignment variables** that synchronize on binding
 - **Threads** that define a sequence of executing instructions
- By design, it has **no observable nondeterminism** (no race conditions)
 - Deterministic dataflow is a form of functional programming
 - « **Concurrency for Dummies** » is the best concurrent paradigm

106


Historical note: concurrency *must* get simpler



- Parallel programming has finally arrived (a surprise to old timers like me!)
 - **Multicore processors**: quad and octo today, 16 and 24 soon, a hundred in a decade, many apps take advantage of it
 - **Distributed computing**: data-intensive with tens of nodes today (NoSQL, MapReduce), hundreds and thousands tomorrow, most apps will do it
- Something fundamental will have to change
 - Sequential programming can't be the default (it's a centralized bottleneck)
 - Libraries cannot hide the issue (interface complexity, distribution structure)
- Concurrency **must become easy**
 - **"Concurrency for Dummies" is the right paradigm** (deterministic dataflow)!
 - It can be used easily on multicore and distributed systems
 - High performance becomes easy
 - Network transparency (program code is the same for different numbers of cores)
 - Modular fault tolerance is easy

107

But is determinism the right default? Yes!



A client/server can't be written in a deterministic paradigm!

It's because the server must accept requests nondeterministically from the two clients

- Deterministic dataflow has strong limitations!
 - A program that needs nondeterminism can't be written
 - Even a simple **client/server can't be written**
- But determinism has enormous advantages, so it is the correct default
 - **Race conditions are impossible** by design
 - With determinism as default, we can **reduce the need for nondeterminism** (in the client/server, it's needed only at the point where the server accepts requests)
 - **Any functional program can be made concurrent** without changing the result

Not a problem!
Just add nondeterminism exactly where it is needed

108

History of deterministic dataflow



- **Deterministic concurrency** has a long history that starts in 1974
 - Gilles Kahn. The semantics of a simple language for parallel programming. In *IFIP Congress*, pp. 471-475, 1974. *Deterministic concurrency*.
 - Gilles Kahn and David B. MacQueen. Coroutines and networks of parallel processes. In *IFIP Congress*, pp. 993-998, 1977. *Lazy deterministic concurrency*.
- Why was it forgotten for so long?
 - Message passing and monitors arrived at about the same time:
 - Carl Hewitt, Peter Bishop, and Richard Steiger. A universal modular ACTOR formalism for artificial intelligence. In *3rd International Joint Conference on Artificial Intelligence (IJCAI)*, pp. 235-245, Aug. 1973.
 - Charles Antony Richard Hoare. Monitors: An operating system structuring concept. *Communications of the ACM*, 17(10):549-557, Oct. 1974.
 - **Actors and monitors express nondeterminism, so they are better. Right?**
- **Dataflow computing** also has a long history that starts in 1974
 - Jack B. Dennis. First version of a data flow procedure language. *Springer Lecture Notes in Computer Science*, vol. 19, pp. 362-376, 1974.
 - **Dataflow remained a fringe subject since it was always focused on parallel programming,** which only became mainstream with the arrival of multicore processors in mainstream computing (e.g., IBM POWER4, the first dual-core processor, in 2001).