

Introduction to Resilient Concurrent/Distributed Programming in Erlang

LINFO1104 – LSINC1104
2023

Peter Van Roy
Université catholique de Louvain

This lecture is based on information taken from many Erlang documents

1

Overview

- Erlang performance
 - Some numbers to show Erlang's strengths
- Basic concepts of the Erlang language
 - Pure functional core with dynamic typing
 - Multi-agent actors: processes (agents) and message passing with mailboxes
 - Failure detection with linking and monitoring
 - Dynamic code updating
- Programming abstractions of the Erlang/OTP platform
 - Standard behaviors (concurrency patterns)
 - Principles of supervisor trees
- Conclusions

2

Erlang strengths

- Multi-agent actor programming
 - Lightweight agents ("Erlang processes") that send each other messages
- Concurrency-oriented programming
 - Support for highly concurrent programming (thousands of processes)
 - Sophisticated libraries to make concurrency and fault tolerance easy ("behavior")
- Distributed programming
 - Processes can be on different computing nodes
- Fault-tolerant applications
 - Support for highly robust software based on "let it fail" philosophy
 - Supervisor trees make fault tolerance easy

3

Erlang introduction


- Erlang was developed in Ericsson for telecommunications in 1986 (Java is from 1991)
 - It is released as **OTP (Open Telecom Platform)** with a full set of libraries
 - It is supported by Ericsson, the Erlang Ecosystem Foundation, and a user community (www.erlang.org)
- Erlang programs consist of lightweight "processes", which are **active agents** that communicate using **asynchronous FIFO message passing** ("actor model" as proposed by Carl Hewitt in 1973)
 - Erlang processes **share nothing**: all data is copied between them
 - Erlang processes receive messages through a mailbox that is accessed by **pattern matching**. Messages can be **received out of order** if they match.
- Erlang/OTP supports reliable long-lived distributed systems using **behaviors** and **supervisors**
 - Behaviors are generic concurrency patterns that make it easy to write robust concurrent systems
 - Supervisors are a general pattern to observe processes and restart them when they crash
 - Ericsson AXD 301 ATM switch with 1.7 million lines of Erlang claims 99.9999999% availability (one may doubt the number of 9's, but the system is extremely available!)

4

Erlang performance

5

Erlang process creation times




From Joe Armstrong, Concurrency Oriented Programming in Erlang, Nov. 2002.

- We compare process creation times for Erlang, Java, and C#
- These numbers were measured in 2002

6

Erlang message sending times




From Joe Armstrong, Concurrency Oriented Programming in Erlang, Nov. 2002.

- We compare message sending times for Erlang, Java, and C#
- These numbers were measured in 2002

7

Use case: web server




From Joe Armstrong, Concurrency Oriented Programming in Erlang, Nov. 2002.

- Throughput versus number of processes for Web servers
 - Red = yaws (Yet Another Web Server, in Erlang on NFS)
 - Green = apache (local disk)
 - Blue = apache (NFS)
- Yaws: 800 KB/s up to 80,000 processes
- Apache: crashes at around 4,000 processes

8

Use case: AXD301 Erlang-based ATM switch

Call Handling Throughput for one CP - AXD 301 release 3.2
Traffic Case: ATM SVC UNI to UNI



From Ulf Wiger, Four-fold Increase in Productivity and Quality, 2001.

Abbreviations:
 CP Control Processor
 ATM Asynchronous Transfer Mode
 SVC Switched Virtual (ATM) Channel
 UNI User Network Interface signaling protocol

- The AXD 301 is a general-purpose high-performance ATM switch from Ericsson
 - The AXD 301 is built using Erlang OTP supplemented with C and Java
- Throughput drops linearly when overloaded
 - 95% throughput at 150% load, descending to 40% throughput at 1000% sustained load
- AXD 301 release 3.2 has 1MLOC Erlang, 900KLOC C/C++, 13KLOC Java
 - In addition, Erlang/OTP at that time had 240KLOC Erlang, 460KLOC C/C++, 15KLOC Java

9

Basic concepts

10

Pure functional core

- Inside each process, Erlang runs as a **pure functional language**
 - All variables are single assignment (bound when they are declared)
 - Functions are values with lexically scoped higher-order programming
 - Pattern matching can be used in **case** and **if** statements (and **receive**)
- All data structures are **symbolic values** (similar to Oz)
 - Integers (arbitrary precision), floats, atoms (symbolic constants)
 - Lists [george,paul,john,ringo] and tuples {Key,Val,L,R}
 - Strings are lists of ASCII codes (integers)
 - **Binary vectors** (used for protocol computations)

11

An Erlang module

- The source code of an Erlang program is organized in **modules**:

```
-module(math).
-export([areas/1]).
-import(lists, [map/2]).

areas(L) -> lists:sum(map(fun(I) -> area(I) end, L)).

area({square,X}) -> X*X;
area({rectangle,X,Y}) -> X*Y.
```

- Modules import and export, giving a **dependency graph of modules**

12

Creating processes and sending messages

- Any process can create another by calling spawn
 - `Pid = spawn(Fun)` : function Fun defines the behavior, Pid is the process name
 - Fun may be anonymous or named
 - `fun(args) -> expr end`
 - `fun name/arity`
- The process name Pid is a unique constant that identifies the process
- Messages can be sent to the process using the process name
 - `Pid ! Message`
 - Messages are sent asynchronously and all data in messages is copied
 - Messages can be received by the `receive` statement

13

Receiving messages

- Each process has a `mailbox` that contains an ordered list of messages received by the process
- Messages are extracted from the mailbox using the `receive` statement
 - The `receive` uses `pattern matching` to remove the first message that matches
 - `receive`

```

pattern1 when guard1 -> expr1;
pattern2 when guard2 -> expr2;
...
patternN when guardN -> exprN
end

```

14

Send and receive

```

Pid ! Message,
...

receive
  Message1 ->
    Actions1;
  Message2 ->
    Actions2;
  ...
  after Time ->
    TimeOutActions
end

```

15

Receive mailbox semantics

- When a process executes **receive**:
 - If the mailbox is empty, the **receive** blocks and waits for a message
 - If the mailbox is not empty, it takes the first message and tries patterns in order starting from the first, if it finds a matching pattern it executes the corresponding code
 - If no pattern matches, the **receive** blocks and waits for the next message
 - Unmatched messages remain in the mailbox and can be removed by future **receive** calls
 - This allows different parts of a process to treat different kinds of messages
 - Messages can be removed out-of-order (in a different order from when they arrived)
 - Care must be taken that messages do not stay in the mailbox forever (memory leak)
- Patterns are symbolic data structures containing variable identifiers and guards are simple built-in tests

16

Process linking

- Two processes can be linked together
 - Process Pid1 calls `link(Pid2)` or conversely; linking is bidirectional
- Process termination: send exit signal
 - A process terminates with an `exit reason`, which is sent as a signal to all linked processes
 - When a process terminates normally, the exit reason is the atom `normal`, otherwise when there is a run-time error, the exit reason is `{Reason,Stack}`
- Propagating process termination: transitive by default ("link set")
 - The default behavior when a process receives an exit signal with reason other than normal is to terminate and to send exit signals with the same reason to its linked processes
 - A process can be set to trap exit signals by calling `process_flag(trap_exit,true)`
 - A received exit signal is then transformed into the message `{'EXIT', FromPid, Reason}`, which is put into the mailbox of the process

17

Using process linking

- If a process throws an exception that is not caught at the top level, then the process terminates and broadcasts its exit signal to all linked processes
- A few additional operations are defined to manage this:
 - `exit(Pid,Why)`: send an exit signal to Pid without terminating
 - `exit(Pid,kill)`: send an unstoppable exit to process Pid

```
start() -> spawn(fun go/0).

go() ->
  process_flag(trap_exit, true),
  loop().

loop() ->
  receive
  {'EXIT',Pid,Why} -> ...
  ... -> ..., loop()
end.
```

18

Dynamic code change (1)

- In a real-time system, we would like to change the code without stopping the system
 - Some systems are never supposed to be stopped, e.g., the X2000 satellite control system developed by NASA
 - Hot code changing is difficult in a monolithic programming system, however, Erlang makes it possible because processes are independent (no sharing)
- Erlang allows for **each module to have two versions of its code**
 - All new processes will be dynamically linked to the latest version
 - If the code is changed, then processes can choose to continue with the old code or to use the new code
 - The choice is determined by how the code is called

19

Dynamic code change (2)

- Call the new version (if available)
- Keep calling the old version:

```
-module(m).
loop(Data, F) ->
  receive
    {From,Q} ->
      {Reply,Data1}=F(Q,Data),
      m:loop(Data1, F)
  end.
```

Use new version

```
-module(m).
loop(Data, F) ->
  receive
    {From,Q} ->
      {Reply,Data1}=F(Q,Data),
      loop(Data1, F)
  end.
```

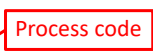
Use old version

- This mechanism is used by libraries that manage upgrading of application releases

20

Client/server process with hot code swap

- Dynamic code change can also be done for **individual processes** using higher-order functions



```

server(Fun, Data) ->
  receive
    {new_fun, Fun1} ->
      server(Fun1, Data);
    {rpc, From, ReplyAs, Q} ->
      {Reply, Data1} =
        Fun{Q, Data},
        From!{ReplyAs, Reply},
        server(Fun, Data1)
  end.

rpc(A, B) ->
  Tag = new_ref(),
  A!{rpc, self(), Tag, B},
  receive
    {Tag, Val} -> Val
  end.

```

21

Programming abstractions for robust software

22

Erlang philosophy

- How can we make robust software?
 - Popular languages (e.g., Java and Python) are inadequate
- Principles of robust software (from Joe Armstrong's Ph.D. thesis)
 - **Errors cannot be fully eliminated**, therefore they must be handled (both hardware and software errors)
 - **Software components are the units of failure**: errors occurring in one will not affect others ("strong isolation")
 - **Software should be fail-fast**: either function correctly or stop quickly
 - **Failure should be detectable** by remote components
 - **Software components share no state**, but communicate through messages

23

Erlang "slogans"

- **"Let it fail"**, "If you can't do your job, crash"
 - Instead of trying to fix things when errors happen, which leads to a large number of complicated states, instead map everything to one simple state, namely "crashed"
- **"Let some other process do error recovery"**
 - Both hardware and software errors can occur and trying to solve them in the process makes things complicated, better to detect and handle any error, hardware or software, elsewhere
- **"Do not program defensively"**
 - Defensive programming means to add checks in the program. This is not productive since it makes the program complicated (in particular, what do you do when a check fails?) and it will not remove all errors. Errors will still occur and still need to be handled. The best way is to map all errors no matter how bizarre to a single fault state, namely "crashed".

24

Erlang/OTP systems

- Erlang/OTP supports a hierarchy of programming abstractions:
 - **Release**: Contains all the information necessary to build and run a system, including a software archive and a set of procedures for installation (including upgrading without stopping)
 - **Application**: Contains all the software necessary to run a single application, not the entire system. Releases are often composed of multiple applications that are largely independent of one another, or that are hierarchically dependent.
 - **Behavior**: A set of processes that together implement a concurrency pattern
 - A notable behavior is **supervisor**: a tree of processes whose purpose is to monitor behaviors and each other and restart them when necessary
 - **Worker**: A process that is an instance of a behavior, usually instances of `gen_server`, `gen_event` or `gen_fsm`

25

Standard Erlang/OTP behaviors

- The Erlang/OTP platform provides **five standard behaviors**:
 - **Generic server (`gen_server`)**: to build client/server architectures with registration, start/stop, timeouts, state management, synchronous/asynchronous calls, error handling
 - **Generic event handler/manager (`gen_event`)**: event handlers, such as loggers, to respond to a stream of events, handling them and sending notifications
 - **Generic finite state machine (`gen_fsm`)**: applications (e.g., protocol stacks) can be modeled as finite state machines, which provides a set of rules $\text{State} \times \text{Event} \rightarrow \text{Actions} \times \text{State}$
 - **Application**: a component that can be started and stopped as a unit, and can be reused in other systems
 - **Supervisor**: the generic toolkit for implementing supervisor hierarchies
- These behaviors hide most of the complexity of each concept, in particularly both concurrency and fault tolerance are vastly simplified using behaviors
 - All non-supervisor behaviors are designed to be pluggable into a supervisor hierarchy

26

Supervisor trees

27

Supervisor trees introduction

- In practical concurrent and distributed systems, it is observed that **most faults and errors are transient**
 - For example: network problems, timing problems, concurrent startup
 - Simple retrying is a surprisingly successful strategy
- Supervisor trees are **designed to favor this strategy**
 - Process sets are “supervised” (observed for failure) by supervisor processes
 - Supervisors have authority to stop and restart supervised processes
 - Supervisors are themselves observed, in case they fail
- Supervisor trees are carefully implemented to avoid races
 - Starting of a supervisor tree is synchronous, to establish correct initial state


28

Supervisor structure and principles

- Supervisor tree is a hierarchy with a root
 - A supervisor tree consists of a set of supervisor nodes, organized as a hierarchy with a root node and internal nodes (the root is a very important process!)
 - A supervisor node is responsible for starting, stopping, and monitoring its child processes
 - All Erlang behaviors are designed to work together with supervisors
- Restart principles
 - Restart strategy: **one_for_one**, **one_for_all**, **rest_for_one**
 - Restart frequency: the number of restarts is limited per time interval
 - If the limit is exceeded, the supervisor terminates and the next higher level supervisor takes some action
 - The intention is to prevent a situation where a process dies repeatedly for the same reason and is always restarted

29

Supervision hierarchies

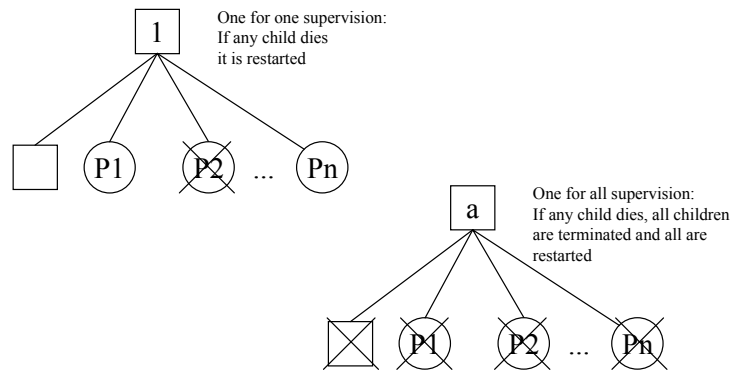


Abbreviations:
 a One-for-all supervision
 1 One-for-one supervision

- A supervisor (□) is a process whose sole purpose is to start, monitor, and possibly restart workers (○)
- A worker is an instance of a behavior, which raise exceptions when errors occur

30

One-for-one and one-for-all supervision



- Different forms of supervision depending on how the children processes work
 - **One-for-one:** if the children are independent (each manages one connection)
 - **One-for-all:** if the children are collaborating, then if one crashes they all have to be restarted, even the correct ones

31

Conclusions

32

Conclusions

- The **Erlang/OTP platform** combines the Erlang language with generic OTP libraries for building resilient highly concurrent and distributed systems (www.erlang.org)
 - **Erlang** supports “processes” (**active agents**) that **share no state** and communicate through **asynchronous messages**
 - Processes are defined using **pure functional programming** and support **dynamic code updating**
 - Failure detection is part of the message communication (**process linking**)
 - **OTP** supports resilient releases and applications using **behaviors**, **supervisor trees**, and **testing**
 - A behavior is a generic concurrency pattern (server, FSM, event handling)
 - A supervisor tree manages how failure handling is done
 - Testing tools for testing concurrent and distributed applications including fault injection
- Erlang/OTP is used successfully for many industrial applications, and Erlang ideas are incorporated into other programming languages and systems
 - Mainstream languages are extended to message passing and agents, and Erlang itself is evolving to resemble mainstream languages (**Elixir** provides a Java-like syntax with an Erlang semantics)
 - The **Erlang Ecosystem Foundation** was founded recently to support Erlang and Elixir

33

Bibliography

34

Bibliography

- Joe Armstrong, [Concurrency Oriented Programming in Erlang](#), Talk slides, Nov. 2002.
- Joe Armstrong, [Making Reliable Distributed Systems in the Presence of Software Errors](#), Ph.D. dissertation, KTH, Dec. 2003.
- Joe Armstrong, [Programming Erlang: Software for a Concurrent World](#), The Pragmatic Bookshelf, 2007.
- Staffan Blau and Jan Rooth. [AXD 301—A New Generation ATM Switching System](#), Ericsson Review No. 1, 1998.
- Francesco Cesarini and Steve Vinoski. [Designing for Scalability with Erlang/OTP](#), O'Reilly, 2016.
- Ericsson AB. [Erlang \(Condensed\)](#) , Talk slides.
- Ericsson AB. [OTP Design Principles](#).
- Ericsson AB. [Erlang/OTP System Documentation Version 10.7](#), March 2020.
- Frederic Trottier-Hebert. [Learn You Some Erlang For Great Good](#).
- Ulf Wiger. [Four-fold Increase in Productivity and Quality](#), March 2001.